


Deep Log-polar networks

Comparison to SITHCon

SITHCon

- Principal operation is log-compression of time series to create $f(t, i)$
 - i indexes τ_{star}
- Learned 1D convolution over τ_{star} , then maxpooling

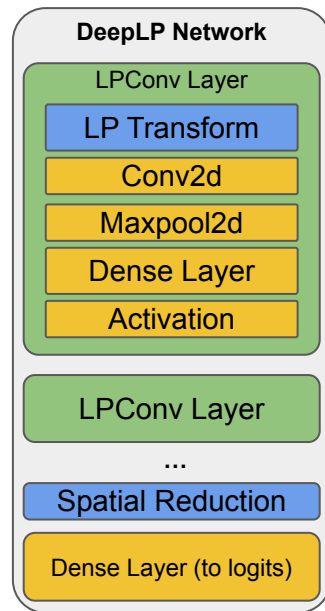



→
single τ_{star} axis

DeepLP

- Principal operation is log-compression of space across rays from a set of angles θ to create $f(x, y, i, j)$
 - i indexes τ_{star}
 - j indexes θ
- Learned 2D convolution over $\tau_{\text{star}} + \theta$, then maxpooling.

many τ_{star} axes






Fig. 2. Rotation and Scaling transformations to a Euclidean image can be read as horizontal and vertical shift respectively, after a log-polar transformation. The log-polar transformation translates rotation and scale in Euclidean images into vertical and horizontal translations (respectively) in the log-polar model.



(figure from Remmelzwaal et al.)

Computing Log-polar transform

Each (τ^*, θ) pair has a corresponding Cartesian coordinate (x'_{ij}, y'_{ij}) :

$$x'_{ij} = \tau_i^* \cos \theta_j = (1 + c)^i \tau_{\min} \cos \theta_j$$

$$y'_{ij} = (1 + c)^i \tau_{\min} \sin \theta_j$$



Get $f(x, y, i, j)$ by:

- Taking the inner product with a window: iSITH in τ^* direction, a Gaussian in the orthogonal direction, with peak at (x'_{ij}, y'_{ij})
- Could also take the average of all pixels in the "pie slice" centered at (x'_{ij}, y'_{ij}) , or use bilinear interpolation to get value at (x'_{ij}, y'_{ij})

LPConv Layer



- Log-polar coordinate transform is computed at every point in the image, resulting in a tensor of size

$H \times W \times C_1 \times N_{\text{tau}} \times N_{\text{theta}}$

- After conv and pooling, the size is

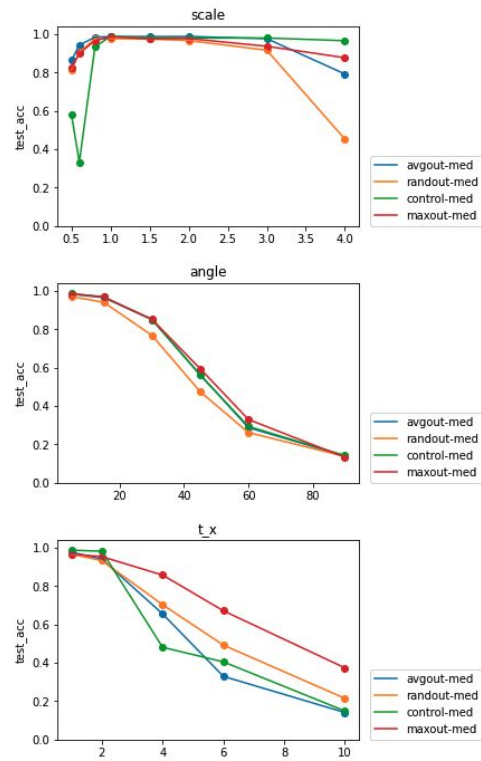
$H \times W \times C_2 \times 1 \times 1$

Computation of log-polar feature map at a single pixel (x, y)



(figure from Su et al., a similar approach but without convolution or max-pooling)

Results on MNIST and RotMNIST




Train small, models (13.5K params) on MNIST, test on rot/scaled/translated images. Comparison of different ways to reduce spatial dimension in last layer (control = choose center pixel). Full max-pooling over only tau dimension.

Note: The rotation-invariant tests are with small ntau and tau_max (to reduce on training time). Scale invariance suffers, though.

Using two models, with 148k and 81k parameters, respectively:

- Training and testing on both RotMNIST gives **98.5%** accuracy for the large model and **97.3%** for the medium model
- Training on only MNIST, then testing on RotMNIST gives **97.8%** and **97.4%** accuracy



rotation invariance
out-of-the-box from
training on
non-augmented MNIST